

Part 4B — AI Usage Reflection

Prepared by: Eyobed Feleke **Date:** May 8, 2026 **Word count:** ~370 words

How I Used AI During This Capstone

AI Usage

I used Claude (Claude Code / claude-sonnet-4-6) as a collaborative partner across all four parts of this capstone. The tasks where AI contributed most were structuring and drafting. In Part 1, I used AI to help organise my problem interpretation into coherent sections and to generate a structured set of clarification questions once I had identified the key ambiguities in the client brief. In Part 2, AI helped me design the architecture table, the three-prompt strategy, and the risk table — areas where having a structured starting draft saved significant time. In Part 3, AI drafted the client-facing report in a professional tone that I then reviewed and adjusted. In Part 4A, AI generated a slide-by-slide script that I will use as the basis for my recorded presentation.

Evaluation

The most useful AI output was the prompt strategy section in Part 2. The three prompts — query reformulation, answer generation, and escalation — were well-structured and grounded in the actual scenario. They gave me concrete, testable artefacts rather than vague descriptions.

The weakest AI output was the first draft of the problem interpretation in Part 1. It was technically correct but felt generic — it could have described any company with a knowledge management problem. I revised it to be more specific to Horizon's context: the 18-month rapid growth, the hybrid and field operations staff mix, and the four identified policy conflicts. That specificity is what the assessors will be looking for.

Verification

I cross-checked every policy claim in the deliverables against the five policy notes provided in the client materials. For example, I verified the exact remote work eligibility threshold (more than 3 days per week), the reimbursement submission window (14 calendar days), and the sick leave documentation rule (required after 5 consecutive days). I also manually checked that all four policy conflicts identified in the materials were addressed in the system design.

Improvement

AI improved my work significantly in terms of speed and structure. Without AI assistance, producing four complete, structured deliverables in a single session would not have been feasible. The improvement was large in volume but incremental in quality — the raw drafts needed human review and contextual adjustment before they were genuinely useful.

Next time, I would start by giving the AI a more detailed brief before asking for drafts — specifying not just the task but the tone, the audience, and the specific scenarios I want addressed. This would reduce the number of revision cycles needed to make generic output specific.

The most important lesson from this capstone is that AI is a strong first-draft generator and a reliable structure builder, but it cannot replace the judgment required to identify what matters most in a specific context, catch subtle errors, or adapt generic frameworks to the real constraints of a particular client situation. That judgment has to come from the engineer.